

RHEL 2-Node High-Availability Clusters for Apache Web Servers

We are configuring 2 node cluster hosting apache webservices.

Node1 & node2 are the clustered node which hosts webservices and belong to a domain called testibm.com

manager node provides storage and also help for the management mode of access for cluster.

Install luci package on the management node, which is used to manage cluster.

Yum install luci ; chkconfig luci on ; service luci restart

```
[root@node4 ~]# service luci restart
Stop luci... [ OK ]
Start luci... [ OK ]
Point your web browser to https://manager.testibm.com:8084 (or equivalent) to access luci
[root@node4 ~]#
```

On node1 & 2 ricci and luci will be installed.

#yum install -y ricci ; service ricci start ; chkconfig ricci on

set password for ricci same in all the nodes.

Access management console for the clusters

https://manager.testibm.com:8084/login?came_from=https%3A%2F%2Fmanager.testibm.com%3A8084%2Fhomebase%2F

Most Visited ▾ Red Hat Customer Portal Documentation Red Hat Network IBM ▾

High Availability management

Homebase

Login

Username

Password

Login

High Availability management

Homebase Clusters

Homebase

Manage Clusters

+ Add * Create ✕

Name

Create New Cluster

Cluster Name

☐ Use the Same Password for All Nodes

Node Name	Password	Ricci Hostname	Ricci Port
<input type="text"/>	<input type="password"/>	<input type="text"/>	<input type="text" value="11111"/> ✕

Add Another Node

☐ Download Packages

☒ Use Locally Installed Packages

☐ Reboot Nodes Before Joining Cluster

☐ Enable Shared Storage Support

Create Cluster Cancel

Create cluster once logged on to luci web interface.

We'll get Create New Cluster called webcluster.

Fill all the mandatory fields as below:

Create New Cluster

Cluster Name

webclust

☒ Use the Same Password for All Nodes

Node Name	Password	Ricci Hostname	Ricci Port
node1.testibm.com	node1.testibm.com	11111
node2.testibm.com	node2.testibm.com	11111

Add Another Node

☒ Download Packages

☐ Use Locally Installed Packages

☐ Reboot Nodes Before Joining Cluster

☒ Enable Shared Storage Support

Create ClusterCancel

Install all the required packages for cluster nodes.

Homebase

Manage Clusters

Nodes

Fence Devices

Failover Domains

Resources

Service Groups

Configure

+ Add

↺ Reboot

🔗 Join Cluster

🔗 Leave Cluster

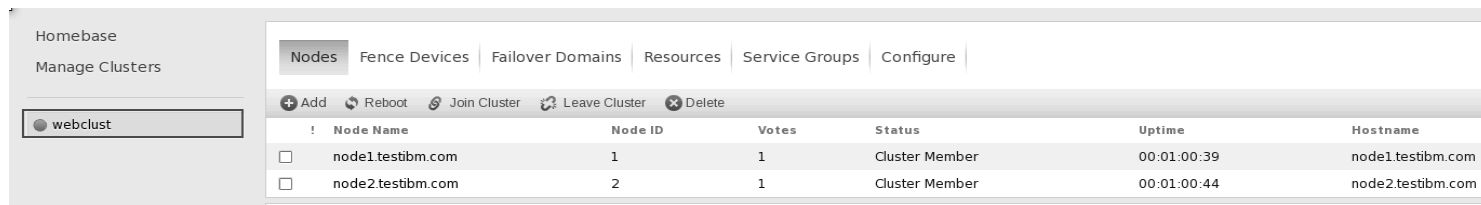
✖ Delete

! Node Name

Please wait...

Creating node "node1.testibm.com" for cluster "webclust": installing packages
Creating node "node2.testibm.com" for cluster "webclust": installing packages

After Packages are installed and Cluster status can be seen as given below.



The screenshot shows the Red Hat Cluster Manager web interface. On the left, there's a sidebar with 'Homebase' and 'Manage Clusters'. Under 'Manage Clusters', a box labeled 'webclust' is selected. The main area has tabs for 'Nodes', 'Fence Devices', 'Failover Domains', 'Resources', 'Service Groups', and 'Configure'. The 'Nodes' tab is active, showing a table of cluster members.

	Node Name	Node ID	Votes	Status	Uptime	Hostname
<input type="checkbox"/>	node1.testibm.com	1	1	Cluster Member	00:01:00:39	node1.testibm.com
<input type="checkbox"/>	node2.testibm.com	2	1	Cluster Member	00:01:00:44	node2.testibm.com

From a cluster node execute clustat which shows the cluster status.

```
[root@node1 ~]# clustat
Cluster Status for webclust @ Mon Jul 21 11:43:27 2014
Member Status: Quorate

Member Name                                ID    Status
-----
node1.testibm.com                          1    Online, Local
node2.testibm.com                          2    Online

[root@node1 ~]#
```

Cluster is up and running, web services has to be hosted on it.

On management node create a storage target which can cat as shareds storage pool.

Create an LVM of 4GB and install a package called scsi-target utils. Also modify the /etc/tgt.target.conf file which should look like given below where in we're using two different network to provide the storage high availability 172.18.1.0 & 172.19.1.0 so that clients access the storage on both the networks.

```
[root@node4 ~]# tail -8 /etc/tgt/targets.conf
<target iqn.2012.04.com.example:cluster1:clusterstorage>
backing-store    /dev/vgsrv/clusterstorage
vendor_id        rhtraining
scsi_sn          rh436
initiator-address 172.18.1.0/24
initiator-address 172.19.1.0/24
</target>

[root@node4 ~]#
```

After modifying the target.conf restart/start tdt daemon and make sure daemon is started after every reboot.

```
[root@node4 ~]# service tgtd start
[root@node4 ~]# service tgtd restart
Stopping SCSI target daemon: [ OK ]
Starting SCSI target daemon: [ OK ]
[root@node4 ~]# chkconfig tgtd on
[root@node4 ~]#
```

Execute tgt-admin -s to see the status.

```
[root@node4 ~]# tgt-admin -s
Target 1: ign.2012.04.com.example:cluster1:clusterstorage
System information:
  Driver: iscsi
  State: ready
I_T nexus information:
LUN information:
  LUN: 0
    Type: controller
    SCSI ID: IET      00010000
    SCSI SN: beaf10
    Size: 0 MB, Block size: 1
    Online: Yes
    Removable media: No
    Readonly: No
    Backing store type: null
    Backing store path: None
    Backing store flags:
  LUN: 1
    Type: disk
    SCSI ID: IET      00010001
    SCSI SN: rh436
    Size: 4295 MB, Block size: 512
    Online: Yes
    Removable media: No
    Readonly: No
    Backing store type: rdwr
    Backing store path: /dev/vgsrv/clusterstorage
    Backing store flags:
Account information:
  ACL information:
    172.18.1.0/24
    172.19.1.0/24
[root@node4 ~]#
```

No iscsi target is configured.

Below are the configuration for initiator to make use of shared storage.

Install initiator package on both the nodes:

```
yum install -y iscsi initiator-utils
```

Discover & login to the storage on both the nodes as below:

```
[root@node1 ~]# iscsiadm -m discovery -t st -p 172.18.1.4
Starting iscsid: [ OK ]
172.18.1.4:3260,1 iqn.2012.04.com.example:cluster1:clusterstorage
[root@node1 ~]# iscsiadm -m discovery -t st -p 172.19.1.4
172.19.1.4:3260,1 iqn.2012.04.com.example:cluster1:clusterstorage
[root@node1 ~]# iscsiadm -m node -T iqn.2012.04.com.example:cluster1:clusterstorage -p 172.18.1.4 -l
Logging in to [iface: default, target: iqn.2012.04.com.example:cluster1:clusterstorage, portal: 172.18.1.4,3260] (multiple)
Login to [iface: default, target: iqn.2012.04.com.example:cluster1:clusterstorage, portal: 172.18.1.4,3260] successful.
[root@node1 ~]# iscsiadm -m node -T iqn.2012.04.com.example:cluster1:clusterstorage -p 172.19.1.4 -l
Logging in to [iface: default, target: iqn.2012.04.com.example:cluster1:clusterstorage, portal: 172.19.1.4,3260] (multiple)
Login to [iface: default, target: iqn.2012.04.com.example:cluster1:clusterstorage, portal: 172.19.1.4,3260] successful.
[root@node1 ~]#
```

```
[root@node2 ~]# iscsiadm -m discovery -t st -p 172.18.1.4
Starting iscsid: [ OK ]
172.18.1.4:3260,1 iqn.2012.04.com.example:cluster1:clusterstorage
[root@node2 ~]# iscsiadm -m discovery -t st -p 172.19.1.4
172.19.1.4:3260,1 iqn.2012.04.com.example:cluster1:clusterstorage
[root@node2 ~]# iscsiadm -m node -T iqn.2012.04.com.example:cluster1:clusterstorage -p 172.18.1.4 -l
Logging in to [iface: default, target: iqn.2012.04.com.example:cluster1:clusterstorage, portal: 172.18.1.4,3260] (multiple)
Login to [iface: default, target: iqn.2012.04.com.example:cluster1:clusterstorage, portal: 172.18.1.4,3260] successful.
[root@node2 ~]# iscsiadm -m node -T iqn.2012.04.com.example:cluster1:clusterstorage -p 172.19.1.4 -l
Logging in to [iface: default, target: iqn.2012.04.com.example:cluster1:clusterstorage, portal: 172.19.1.4,3260] (multiple)
Login to [iface: default, target: iqn.2012.04.com.example:cluster1:clusterstorage, portal: 172.19.1.4,3260] successful.
[root@node2 ~]#
```

Enable multipath.

On both the nodes make sure device-mapper-multipath package is installed.

Yum list device-mapper-multipath

If not please install the package using “yum install -y device-mapper-multipath”

Enable , start and ensure that multipathd is started after every reboot.

```

[root@node1 ~]# mpathconf --enable
[root@node1 ~]# service multipathd restart
Stopping multipathd daemon: [FAILED]
Starting multipathd daemon: [ OK ]
[root@node1 ~]# chkconfig multipathd on
[root@node1 ~]# multipath -r
Jul 21 12:24:45 | /lib/udev/scsi_id exited with 1
Jul 21 12:24:45 | /lib/udev/scsi_id exited with 1
reload: mpatha (1IET 00010001) undef rhtraini,VIRTUAL-DISK
size=4.0G features='0' hwhandler='0' wp=undef
| +- policy='round-robin 0' prio=1 status=undef
| `-- 2:0:0:1 sda 8:0 active ready running
`+- policy='round-robin 0' prio=1 status=undef
  `-- 3:0:0:1 sdb 8:16 active ready running
[root@node1 ~]# █

```

```

[root@node2 ~]# mpathconf --enable
[root@node2 ~]# service multipathd restart
Stopping multipathd daemon: [FAILED]
Starting multipathd daemon: [ OK ]
[root@node2 ~]# chkconfig multipathd on
[root@node2 ~]# multipath -r
Jul 21 12:25:06 | /lib/udev/scsi_id exited with 1
Jul 21 12:25:06 | /lib/udev/scsi_id exited with 1
reload: mpatha (1IET 00010001) undef rhtraini,VIRTUAL-DISK
size=4.0G features='0' hwhandler='0' wp=undef
| +- policy='round-robin 0' prio=1 status=undef
| `-- 2:0:0:1 sda 8:0 active ready running
`+- policy='round-robin 0' prio=1 status=undef
  `-- 3:0:0:1 sdb 8:16 active ready running
[root@node2 ~]# █

```

From above multipath device is mpatha however we want to give user friendly name or an alias so that it can easily identify which storage we're using for cluster storage.

Make the same changes to both the nodes in multipath.conf. And restart the multipathd daemon.

```

[root@node1 ~]# grep -v '#' /etc/multipath.conf

defaults {
    user_friendly_names yes
}
multipaths {
    multipath {
        wwid                "1IET        00010001"
        alias                clusterstorage
    }
}

blacklist {
}
[root@node1 ~]# service multipathd restart
Stopping multipathd daemon:                [ OK ]
Starting multipathd daemon:                [ OK ]
[root@node1 ~]# multipath -r
Jul 21 12:34:00 | /lib/udev/scsi_id exited with 1
Jul 21 12:34:00 | /lib/udev/scsi_id exited with 1
reload: clusterstorage (1IET        00010001) undef rhtraini,VIRTUAL-DISK
size=4.0G features='0' hwhandler='0' wp=undef
|-+- policy='round-robin 0' prio=1 status=undef
|  `-- 2:0:0:1 sda 8:0    active ready running
|-+- policy='round-robin 0' prio=1 status=undef
|  `-- 3:0:0:1 sdb 8:16   active ready running
[root@node1 ~]#

```



```
[root@node1 ~]# fdisk -cu /dev/mapper/clusterstorage

Command (m for help): p

Disk /dev/mapper/clusterstorage: 4294 MB, 4294967296 bytes
255 heads, 63 sectors/track, 522 cylinders, total 8388608 sectors
Units = sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes
Disk identifier: 0xa747b69b

               Device Boot      Start         End      Blocks   Id  System
               -----
Command (m for help): n
Command action
  e   extended
  p   primary partition (1-4)
p
Partition number (1-4): 1
First sector (2048-8388607, default 2048):
Using default value 2048
Last sector, +sectors or +size{K,M,G} (2048-8388607, default 8388607): +2G

Command (m for help): t
Selected partition 1
Hex code (type L to list codes): 8e
Changed system type of partition 1 to 8e (Linux LVM)

Command (m for help): w
The partition table has been altered!

Calling ioctl() to re-read partition table.

WARNING: Re-reading the partition table failed with error 22: Invalid argument.
The kernel still uses the old table. The new table will be used at
the next reboot or after you run partprobe(8) or kpartx(8)
Syncing disks
```

Execute partprobe or kpartx on both the nodes so that partition is reflected on both the nodes

Make sure locking_type = 3 is reflected on all the cluster nodes.

Now create new clusteredvg & lv named as webdata on new partition & format with gfs2 Filesystem on any of the cluster node.

```
[root@node1 ~]# pvcreate /dev/mapper/clusterstoragep1
Writing physical volume data to disk "/dev/mapper/clusterstoragep1"
Physical volume "/dev/mapper/clusterstoragep1" successfully created
[root@node1 ~]# vgcreate -cy clustervg /dev/mapper/clusterstoragep1
Clustered volume group "clustervg" successfully created
```

```
[root@node1 ~]# lvcreate -l 511 -n webdata clustervg
Logical volume "webdata" created
[root@node1 ~]# vgs
VG          #PV #LV #SN Attr   VSize VFree
clustervg   1   1   0 wz--nc 2.00g   0
vgsrv       1   3   0 wz--n- 8.78g 4.69g
[root@node1 ~]# lvs
LV          VG          Attr   LSize   Origin Snap%   Move Log Copy%   Convert
webdata     clustervg -wi-a- 2.00g
home        vgsrv        -wi-ao 256.00m
root        vgsrv        -wi-ao 3.31g
swap        vgsrv        -wi-ao 544.00m
[root@node1 ~]#
```

```
swap        vgsrv        -wi-ao 544.00m
[root@node1 ~]# man mkfs.gfs2
[root@node1 ~]# mkfs.gfs2 -t webclust:webdata -p lock_dlm -j 2 /dev/clustervg/webdata
This will destroy any data on /dev/clustervg/webdata.
It appears to contain: symbolic link to `../dm-5'

Are you sure you want to proceed? [y/n] y

Device:                /dev/clustervg/webdata
Blocksize:             4096
Device Size            2.00 GB (523264 blocks)
Filesystem Size:       2.00 GB (523264 blocks)
Journals:              2
Resource Groups:       8
Locking Protocol:      "lock_dlm"
Lock Table:            "webclust:webdata"
UUID:                 a57dcf7b-e923-451a-44e2-0af947c189b7

[root@node1 ~]#
```

Now that we've GFS2 Filesystem available we're mounting it temporary in any of the node in our case we're mounting it on node1. Copying/Creating the index.html file under /var/www/html and setting the selinux context as selinux is enforcing in our case.

```
[root@node1 ~]# mount /dev/clustervg/webdata /var/www/html/
[root@node1 ~]# echo "testing with GFS2 Filesystem" >/var/www/html/index.h
[root@node1 ~]# restorecon -Rv /var/www/html/
[root@node1 ~]# ls -ldZ /var/www/html/
drwxr-xr-x. root root system_u:object_r:httpd_sys_content_t:s0 /var/www/ht
[root@node1 ~]# ls -lZ /var/www/html/
-rw-r--r--. root root unconfined_u:object_r:httpd_sys_content_t:s0 index.h
[root@node1 ~]# umount /var/www/html/
[root@node1 ~]#
```

Now our FS is ready which has the data/index.html which we're using to publish our website.

Crate a Failover domain:

Homebase
Manage Clusters

NodesFence DevicesFailover DomainsResourcesService GroupsConfigure

webclust

+ Add

✕ Delete

Name	Prioritized	Restricted
<input type="checkbox"/> webfaildomain	✓	✓

webfaildomain

☒ Prioritized

☒ Restricted

☐ No Failback

Order the nodes to which services failover.

Service can run only on nodes specified.

Do not send service back to 1st priority node when it becomes available again.

Update Properties

Services

Members

Member	Priority	
node1.testibm.com	<input checked="" type="checkbox"/>	1
node2.testibm.com	<input checked="" type="checkbox"/>	2

Update Settings

Click on Resources tab and create new resources:

Homebase
Manage Clusters

NodesFence DevicesFailover DomainsResourcesService GroupsConfigure

webclust

+ Add

✕ Delete

Name/IP

Add Resource to Cluster

GFS2

Name

webdata

Mount Point

/var/www/html

Device, FS Label, or UUID

/dev/clustervg/webdata

Filesystem Type

GFS2

Mount Options

Filesystem ID (optional)

Force Unmount

☒

Reboot Host Node if Unmount Fails

☒

Submit

Cancel

Now Create the IP resource

Homebase
Manage Clusters

Nodes | Fence Devices | Failover Domains | **Resources** | Service Groups | Configure

+ Add **✕ Delete**

Name/IP

☐ webdata

webdata

GFS2

Name
Mount Point
Device, FS Label, or UUID
Filesystem Type
Mount Options
Filesystem ID (optional)

Add Resource to Cluster

IP Address

IP Address

IP Address 172.16.1.50
Netmask Bits (optional) 24
Monitor Link ☒
Disable Updates to Static Routes ☒
Number of Seconds to Sleep After Removing an IP Address 10

Submit Cancel

Creating apache resource

Homebase
Manage Clusters

Nodes | Fence Devices | Failover Domains | **Resources** | Service Groups | Configure

+ Add **✕ Delete**

Name/IP

☐ webdata

☐ 172.16.1.50/24

172.16.1.50/24

IP Address

IP Address
Netmask Bits (optional)
Monitor Link
Disable Updates to Static Routes
Number of Seconds to Sleep After Removing an IP Address

Apache

Name webservice
Server Root /etc/httpd
Config File conf/httpd.conf
httpd Options
Shutdown Wait (seconds) 3

Submit Cancel

Apply

All three resources are created.

✓ Created global resource "webservice"

Homebase
Manage Clusters

● webclust

Nodes | Fence Devices | Failover Domains | **Resources** | Service Groups | Configure

+ Add ✕ Delete

Name/IP	Type	In Use
<input type="checkbox"/> webdata	GFS2	No
<input type="checkbox"/> 172.16.1.50/24	IP Address	No
<input type="checkbox"/> webservice	Apache Server	No

webservice

Next step is to create a Service group.

Click on Service group Tab

Nodes | Fence Devices | Failover Domains | Resources | **Service Groups** | Configure

+ Add ▶ Start ↺ Restart ■ Disable ✕ Delete

!	Name	Status	Autosta
No item to display			

Follow the same steps to add the resources in the webhttp service (IP, Filesystem, and web resource /apache)

Add Service Group to Cluster

Service Name

webhttp

Automatically Start This Service



Run Exclusive



Failover Domain

webfaildomain

Recovery Policy

Relocate

Restart Options

Maximum Number of Restart Failures Before Relocating

Length of Time in Seconds After Which to Forget a Restart

Add Resource

Submit

Cancel

[Remove](#)

IP Address

IP Address

Netmask Bits (optional)

Monitor Link



Disable Updates to Static Routes



Number of Seconds to Sleep After Removing an IP Address

Independent Subtree



Non-Critical Resource



Independent Subtree/Non-Critical Options

Maximum Number of Failures

Failure Expire Time (seconds)

Maximum Number of Restarts

Restart Expire Time (seconds)

[Add a child resource](#)

[Add Resource](#)

[Submit](#)

[Cancel](#)

GFS2

Name	<input type="text" value="webdata"/>
Mount Point	<input type="text" value="/var/www/html"/>
Device, FS Label, or UUID	<input type="text" value="/dev/clustervg/webdata"/>
Filesystem Type	<input type="text" value="GFS2"/>
Mount Options	<input type="text"/>
Filesystem ID (optional)	<input type="text" value="42848"/>
Force Unmount	<input checked="" type="checkbox"/>
Reboot Host Node if Unmount Fails	<input checked="" type="checkbox"/>
Independent Subtree	<input type="checkbox"/>
Non-Critical Resource	<input type="checkbox"/>

Independent Subtree/Non-Critical Options

Maximum Number of Failures	<input type="text"/>
Failure Expire Time (seconds)	<input type="text"/>
Maximum Number of Restarts	<input type="text"/>
Restart Expire Time (seconds)	<input type="text"/>

Once all the resources are added click on submit.

Apache

Name

Server Root

Config File

httpd Options

Shutdown Wait (seconds)

Independent Subtree

☐

Non-Critical Resource

☐

Independent Subtree/Non-Critical Options

Maximum Number of Failures

Failure Expire Time (seconds)

Maximum Number of Restarts

Restart Expire Time (seconds)

Add a child resource

Add Resource

Submit

Cancel

Click on submit.



Created cluster service "webhttp"

Homebase

Manage Clusters

webclust

Nodes Fence Devices Failover Domains Resources **Service Groups** Configure

+ Add ▶ Start ↺ Restart ⏻ Disable ✕ Delete

! Name	Status	Autostart	Failover Domain
<input type="checkbox"/> ? webhttp	Unknown	<input checked="" type="checkbox"/>	webfaildomain

Homebase

Manage Clusters

webclust

Nodes

Fence Devices

Failover Domains

Resources

Service Groups

Configure

+

Add

▶

Start

↺

Restart

■

Disable

✕

Delete

<div>!</div> Name	Status	Autostart	Failover Domain
<div><input type="checkbox"/></div> webhttp	Running on node1.testibm.com	<div><input checked="" type="checkbox"/></div>	webfaildomain

Select an item to view details

clustat output:

```
[root@node1 ~]# clustat
Cluster Status for webclust @ Mon Jul 21 14:28:56 2014
Member Status: Quorate

Member Name                               ID   Status
-----
node1.testibm.com                         1 Online, Local, rgmanager
node2.testibm.com                         2 Online, rgmanager

Service Name                               Owner (Last)   State
-----
service:webhttp                          node1.testibm.com started
[root@node1 ~]#
```

Creating fencing Device Select “Fence Virt (Multicast mode)” .

Homebase
Manage Clusters
webclust

Nodes
Fence Devices
Failover Domains
Resources
Service Groups
Configure

Add
Delete

Name	Fence Type	Nodes Using
<div> <div>Add Fence Device (Instance)</div> <div> <div>Fence virt (Multicast Mode)</div> <div> Fence Type Fence xvm Name virtfence </div> <div> Submit Cancel </div> </div> </div>		

Now add the instance for fencing.

Click on node tab

Homebase
Manage Clusters

Nodes

Fence Devices

Failover Domains

Resources

Service Groups

Configure

webclust

+

Add

↺

Reboot

↻

Join Cluster

↻

Leave Cluster

✖

Delete

!	Node Name	Node ID	Votes	Status	Uptime	Hostname
<input type="checkbox"/>	node1.testibm.com	1	1	Cluster Member	00:02:07:03	node1.testibm.com
<input type="checkbox"/>	node2.testibm.com	2	1	Cluster Member		node2.testibm.com

Select an item to view details

then click on node1.testibm.com

Nodes

Fence Devices

Failover Domains

Resources

Service Groups

Configure

+

Add

↺

Reboot

↻

Join Cluster

↻

Leave Cluster

✖

Delete

!	Node Name	Node ID	Votes	Status	Uptime	Hostname
<input type="checkbox"/>	node1.testibm.com	1	1	Cluster Member	00:02:08:43	node1.testibm.com
<input type="checkbox"/>	node2.testibm.com	2	1	Cluster Member		node2.testibm.com

node1.testibm.com

Status Cluster Member

✖

↺

↻

↻

Properties

Number of votes

Update Properties

ricci host

ricci port

Services

webhttp

Failover Domains

webfaildomain

Priority1

Fence Devices

Method

Add Fence Method

Cluster Daemons

Add Fence Method and given any name.

node1.testibm.com
Status Cluster Member

Properties

Number of votes [Update Properties](#)

ricci host

ricci port

Services

webhttp

Failover Domains

	Priority
webfaildomain	1

Fence Devices

Method

test [Remove](#)

Name	Type/Values
Add Fence Instance	

[Add Fence Method](#)

node1.testibm.com
Status Cluster Member

Properties

Number of votes

ricci host

ricci port

Services

webhttp

Failover Domains

[webfaildomain](#)

Fence Devices

Method

test

[Add Fence Instance](#)

[Add Fence Method](#)

Add Fence Device (Instance)

Domain

[Submit](#) [Cancel](#)

Select the Virtfence which we created and Domain “node1” or UUID of the VM (It depends what name is given to the VM in our case it's node1)

Click on submit.

Updating fence settings for node "node1.testibm.com"

HomebaseManage Clusters

webclust

NodesFence DevicesFailover DomainsResourcesService GroupsConfigure

+ Add↻ Reboot⚙ Join Cluster⚙ Leave Cluster🗑 Delete

Node Name	Node ID	Votes	Status	Uptime	Hostname
<input type="checkbox"/> node1.testibm.com	1	1	Cluster Member	00:02:13:39	node1.testibm.com
<input type="checkbox"/> node2.testibm.com	2	1	Cluster Member		node2.testibm.com

node1.testibm.com
Status Cluster Member

Once added you'll see message as given above.

Same setps to be carried on node2.

Properties

Number of votes

1

Update Properties

ricci host

node2.testibm.com

ricci port

11111

Services

Failover Domains

webfaildomain

Priority

Fence Devices

Method

test

Name

Add Fence Instance

Add Fence Method

Add Fence Device (Instance)

virtfence (xvm Virtual Machine Fencing) ⌵

Domain

node2

Submit

Cancel

Remove

Updating fence settings for node "node2.testibm.com"

HomebaseManage Clusters

webclust

NodesFence DevicesFailover DomainsResourcesService GroupsConfigure

+ Add↻ Reboot⚙ Join Cluster⚙ Leave Cluster🗑 Delete

Node Name	Node ID	Votes	Status	Uptime	Hostname
<input type="checkbox"/> node1.testibm.com	1	1	Cluster Member	00:02:16:30	node1.testibm.com
<input type="checkbox"/> node2.testibm.com	2	1	Cluster Member	00:02:16:27	node2.testibm.com

node2.testibm.com
Status Cluster Member

Properties

Number of votes

1

Update Properties

ricci host

node2.testibm.com

ricci port

11111

On the base machine install the fencing package.

Yum install fence-virtfd fence-virtfd-libvirt fence-virtfd-multicast

Create a fence key as below:

```
#dd if=/dev/urandom of=/etc/cluster/fence_xvm.key bs=4 count=1
```

```
# fence_virtfd -c
```

```
[root@oc0405105284 ~]# fence_virtfd -c
Module search path [/usr/lib64/fence-virt]:

Available backends:
    libvirt 0.1
    libvirt-qpid 0.1
Available listeners:
    multicast 1.1
    serial 0.4

Listener modules are responsible for accepting requests
from fencing clients.

Listener module [multicast]:

The multicast listener module is designed for use environments
where the guests and hosts may communicate over a network using
multicast.

The multicast address is the address that a client will use to
send fencing requests to fence_virtfd.

Multicast IP Address [225.0.0.12]:
```

```
Using ipv4 as family.

Multicast IP Port [1229]:

Setting a preferred interface causes fence_virtfd to listen only
on that interface. Normally, it listens on the default network
interface. In environments where the virtual machines are
using the host machine as a gateway, this *must* be set
(typically to virbr0).
Set to 'none' for no interface.

Interface [virbr1]:
```

Specify the Interface which is being used to communicate to the VMs from the (physical) base machine.

The key file is the shared key information which is used to authenticate fencing requests. The contents of this file must be distributed to each physical host and virtual machine within a cluster.

Key File [/etc/cluster/fence_xvm.key]:

Backend modules are responsible for routing requests to the appropriate hypervisor or management layer.

Backend module [libvirt]:

The libvirt backend module is designed for single desktops or servers. Do not use in environments where virtual machines may be migrated between hosts.

Libvirt URI [qemu:///system]:

Configuration complete.

=== Begin Configuration ===

```
backends {
    libvirt {
        uri = "qemu:///system";
    }
}

listeners {
    multicast {
        interface = "virbr1";
        port = "1229";
        family = "ipv4";
        address = "225.0.0.12";
        key_file = "/etc/cluster/fence_xvm.key";
    }
}
```

```
fence_virt {
    module_path = "/usr/lib64/fence-virt";
    backend = "libvirt";
    listener = "multicast";
}
```

=== End Configuration ===

Replace /etc/fence_virt.conf with the above [y/N]? y
[root@oc0405105284 ~]#

now copy the /etc/cluster/fence_xvm.key to all the cluster nodes.

Make sure fence_virt daemon is running on base machine.

/etc/init.d/fence_virttd status

fence_virttd (pid 11575) is running...

Then login to cluster node and test it.

Below command can be used to see fencing is working or not.

fence_xvm -o list

To check if service is getting failover as expected :

```
[root@node1 ~]# clusvcadm -r webhttp
Trying to relocate service:webhttp...Success
service:webhttp is now running on node2.testibm.com
[root@node1 ~]# clustat
Cluster Status for webclust @ Mon Jul 21 16:00:21 2014
Member Status: Quorate

Member Name                                ID    Stat
-----
node1.testibm.com                          1    Onl
node2.testibm.com                          2    Onl

Service Name                                Owner (L
-----
service:webhttp                            node2.te
[root@node1 ~]#
```

Service has been started on node2.

```
[root@node2 ~]# clustat
Cluster Status for webclust @ Mon Jul 21 16:06:26 2014
Member Status: Quorate

Member Name                                ID      Status
-----
node1.testibm.com                          1 Online, rgmanager
node2.testibm.com                          2 Online, Local, rgmanager

Service Name                               Owner (Last)                               State
-----
service:webhttp                            node2.testibm.com                           start
[root@node2 ~]#
```

Make sure if all the services are on at startup so it starts after reboot.

```
[root@node1 ~]# chkconfig --list |egrep -i "cman|clvmd|gfs2|rgmanager|ricc
clvmd          0:off  1:off  2:on   3:on   4:on   5:on   6:off
cman           0:off  1:off  2:on   3:on   4:on   5:on   6:off
gfs2           0:off  1:off  2:on   3:on   4:on   5:on   6:off
modclusterd   0:off  1:off  2:on   3:on   4:on   5:on   6:off
rgmanager      0:off  1:off  2:on   3:on   4:on   5:on   6:off
ricci         0:off  1:off  2:on   3:on   4:on   5:on   6:off
[root@node1 ~]#
```

```
[root@node2 ~]# chkconfig --list |egrep -i "cman|clvmd|gfs2|rgmanager|ricc
clvmd          0:off  1:off  2:on   3:on   4:on   5:on   6:off
cman           0:off  1:off  2:on   3:on   4:on   5:on   6:off
gfs2           0:off  1:off  2:on   3:on   4:on   5:on   6:off
modclusterd   0:off  1:off  2:on   3:on   4:on   5:on   6:off
rgmanager      0:off  1:off  2:on   3:on   4:on   5:on   6:off
ricci         0:off  1:off  2:on   3:on   4:on   5:on   6:off
[root@node2 ~]#
```

Testing fencing as we'll failover.

```
[root@node1 ~]# fence_xvm -o list
node1          16ae4af9-6416-f008-123f-1c1853ef7666 on
node2          be02ea85-1985-50ee-0671-2e5a6dbf0b0b on
node4          dd2c9519-1111-b377-f644-9fe7eb5ce97d on
[root@node1 ~]# fence_node node2
fence node2 success
[root@node1 ~]#
```

```
[root@node1 ~]# clustat
Cluster Status for webclust @ Mon Jul 21 16:13:33 2014
Member Status: Quorate
```

Member Name	ID	Status
node1.testibm.com	1	Online, Local, rgmanager
node2.testibm.com	2	Offline

Service Name	Owner (Last)	State
service:webhttp	node1.testibm.com	started

```
[root@node1 ~]#
```